

15-712 Project Report

Introduce Indirection Scheduling to Datacenter Hybrid Switch

Conglong Li, Anbang Zhao
{conglonl, anbangz}@andrew.cmu.edu

1 Introduction

Traditionally, circuit switching and packet switching have been considered as alternative design choices, each with their own advantages and disadvantages. Packet switches are efficient at multiplexing traffic across a large number of ports, while circuit switches can often service higher line rates in a more cost-effective manner, especially when combined with optical links. However, as optical circuit switching technology advances, the reconfiguration time (i.e., the amount of time it takes to alter the input-to-output mapping) is swiftly decreasing, thus increasingly blurring the division between the packet and circuit regimes.

One key domain in which these design tradeoffs are being revisited today is datacenter switching. The datacenter environment, aggregating tremendous amounts of compute and storage capacity, has driven demand for ever increasing port counts and line speeds. However, supporting these demands with existing packet switching technology is becoming increasingly expensive—in cost, heat and power. Moreover, studies have shown that the traffic demand inside a datacenter (both within and across racks) is frequently highly nonuniform, with a large fraction of the traffic at each switch port destined for a small number of output ports [6].

Based on these observations, researchers had proposed hybrid datacenter network architectures that offer higher throughputs at lower cost by combining switching technologies. In particular, recent proposals suggest employing highspeed optical [2, 3, 9] or wireless [4, 5, 10] networks configured to service the heavy flows, while passing the remainder of the traffic through a traditional, relatively underprovisioned packet-switched network.

As technology trends usher in dramatically faster reconfiguration times, the distinction between packet and circuit is blurred, and ever smaller flows can take advantage of a hybrid fabric. This trend will soon allow servicing the bulk of the traffic through a rapidly reconfigurable optical switch [8], leaving a relatively minor portion to be serviced by the packet network [7]. While the potential cost savings that hybrid technologies could realize is large, the design space for scheduling resources in the hybrid regime is not yet well understood.

Many of the classical approaches to scheduling for switches with non-trivial reconfiguration delays divide the offered demand into two parts: an initial, heavy-weight component that is served by $O(N)$ highly utilized configurations with significant durations, and a second, residual component that is serviced by a similar number of short, under-utilized schedules. Recent Solstice scheduling

algorithm [7] exploits the skewed nature of datacenter traffic patterns to create a small number of configurations with long durations that minimize the penalty for reconfiguration and leaves only a small amount of residual demand to be serviced by a low-speed (and lowcost) unconstrained packet switch.

To ensure high overall circuit utilization, each circuit configuration must be maintained for a relatively long period with respect to the reconfiguration delay. This leads to two issues for the scheduling algorithm: Since the demand is skewed, the demands served by a single configuration could have a large variation. Thus a portion of the links in optical switch could be under-utilized. Since the configuration time is still non-trivial, it is still hard to schedule small traffic demands efficiently on optical switch. Scheduling algorithms are forced to serve these small demands by packet switch, which is slower and less cost-effective.

To increase the utilization rate of the optical switches and to reduce the number of configurations, we propose an indirection scheduling technique: instead of scheduling each demand with unique direct path, we could leverage existing paths to serve some demands indirectly. Using this technique, scheduling algorithm could reduce the skewness of the demands and reduce the number of configurations.

The evaluation of the indirection technique has two objectives. The first is to find out the optimal possible benefit of indirection. We achieved this goal by formulating the scheduling problem into an integer linear programming (ILP) problem. Using *Gurobi* ILP solver [1] to compare the optimal solution with and without indirection, we show that using indirection could reduce the number of configurations and the total transmission time.

The second objective is to find out the practical benefit of indirection. We achieved this goal by simulations of hybrid switch scheduling. We implemented two simple indirection heuristics and applied them to the Solstice algorithm. Comparing the simulations of Solstice algorithm with and without indirection heuristics, results show that even simple indirection heuristics could provide improvements above direct path algorithm. Although our indirection heuristics didn't approach similar amount of improvement as in the ILP formulation results, we argue that better improvement is possible with clever indirection heuristics embedded with the scheduling algorithm.

The rest of this report is organized as follows. Section 2 described the ILP formulation and the solver results. Section 3 describes the preliminary indirection heuristics and the simulation results. Finally, we summarize this project in Section 4.

2 ILP Formulation of the Scheduling Problem

2.1 Problem Formulation

Integer linear programming is a mathematical optimization method. An integer linear program is expressed with a minimization/maximization objective function, constraints of the variables, and bounds of the variables. For our scheduling problem, the objective function is clearly the minimization of the total transmission time, but the constraints are complicated. In the remaining of this section, we will describe the constraints implemented in the program.

Input Demand Matrix To define the input of each integer linear program, we defined the demands needs to be served as a N by N matrix, where N is the number of ports in the network. Entry (a, b) in the input demand matrix denotes the amount of demand to be transfered from port a to port b . Each (a, a) entry has zero demand. Then we ask the program to solve the scheduling problem so

Optimal Total Time	Random $\delta = 5$	Random $\delta = 10$	Random $\delta = 20$	Skewed $\delta = 10$
Orig	100%	100%	100%	100%
Indi	93.93%	90.65%	83.8%	97.9%

Table 1: Normalized average optimal solution achieved by direct path scheduling (Orig) and indirection technique (Indi), δ is the reconfiguration time.




Figure 1: Optimal solution count for different number of configurations with random demand




Figure 2: Optimal solution count for different number of configurations with skewed demand

that each demand in the input demand matrix must be served by either optical or packet switch.

Optical Switch Constraints To define the optical switch, we need to implement the concept of configurations. Each configuration is a matrix that denotes the demands served by this configuration. Since the input-output mapping is fixed in a single configuration, each row and each column can only have no more than one non-zero entry. For each new configuration scheduled, an additional reconfiguration time is added to the total transmission time.

Packet Switch Constraints Packet switch constraints are different. Packet switch doesn't have the concept of configurations. It could schedule arbitrary amount of demands on each link without any reconfiguration time. However, packet switch has less capacity than optical switch. Thus it could only schedule a smaller fraction of demands within the same time duration. In our constraints, we configure the packet switch to have 1/10 capacity of the optical switch.

Indirection Technique Constraints To define the indirection technique, we defined each possible indirection decision in the program. If a decision is made that certain amount of demand from port a to port b will be indirectly served by paths a -> c and c -> b, the input demand matrix will reduce the demand in entry (a, b) and increase the demands in entries (a, c) and (c, b).

2.2 ILP Solving Results

We used the Gurobi ILP solver to solve the optimal solutions of our integer linear programs. Since integer linear solving is NP-hard, we choose to solve 5 by 5 input demand matrix. This is a small input compared to the matrix used in the simulator. But it is big enough to show the potential benefit of the indirection technique.

We used two different kinds of workload. The first is random workload where each entry has a demand randomly chosen from 0 to 40. To show the effect of the reconfiguration time, we tested different reconfiguration times between 5 to 20 for the random workload. The second workload is skewed workload where each port sends one big demand (around 250) and two small demand (around 40). For this workload, we only tested with reconfiguration time as 10.

We generated 20 integer linear programs with different input demand matrices for each workload and each reconfiguration time used. Then we let the Gurobi solver to solve these programs and report the optimal solutions. Table 1 shows the normalized average optimal solutions for each workload and different reconfiguration time. Indirection technique provides improvement for all workloads. In particular, indirection provides more improvements when the ratio of demand to reconfiguration time is smaller. This is because indirection technique helps on reducing the number of configurations. Figure 1 and Figure 2 depict the number of configurations needed for each optimal solution for both with and without indirection. Results show that indirection technique reduced the number of configurations for most of test cases.

3 Simulation Results

4 Conclusion

References

- [1] Gurobi Optimizer. URL <http://www.gurobi.com/>.
- [2] K. Chen, A. Singla, A. Singh, K. Ramachandran, L. Xu, Y. Zhang, X. Wen, and Y. Chen. OSA: An optical switching architecture for data center networks with unprecedented flexibility. In *USENIX NSDI*, 2012.
- [3] N. Farrington, G. Porter, S. Radhakrishnan, H. H. Bazzaz, V. Subramanya, Y. Fainman, G. Papen, and A. Vahdat. Helios: a hybrid electrical/optical switch architecture for modular data centers. In *SIGCOMM 2010*.
- [4] D. Halperin, S. Kandula, J. Padhye, P. Bahl, and D. Wetherall. Augmenting data center networks with multi-gigabit wireless links. In *SIGCOMM 2011*.
- [5] S. Kandula, J. Padhye, and P. Bahl. Flyways to de-congest data center networks. In *HotNets 2009*, .
- [6] S. Kandula, S. Sengupta, A. Greenberg, P. Patel, and R. Chaiken. The nature of data center traffic: measurements & analysis. In *IMC 2009*, .
- [7] H. Liu, F. Lu, A. Forencich, R. Kapoor, M. Tewari, G. M. Voelker, G. Papen, A. C. Snoeren, and G. Porter. Circuit Switching Under the Radar with REACToR. In *USENIX NSDI*, 2014.
- [8] G. Porter, R. Strong, N. Farrington, A. Forencich, P. Chen-Sun, T. Rosing, Y. Fainman, G. Papen, and A. Vahdat. Integrating microsecond circuit switching into the data center. In *SIGCOMM 2013*.
- [9] G. Wang, D. G. Andersen, M. Kaminsky, K. Papagiannaki, T. Ng, M. Kozuch, and M. Ryan. c-Through: Part-time optics in data centers. In *SIGCOMM 2010*.

- [10] X. Zhou, Z. Zhang, Y. Zhu, Y. Li, S. Kumar, A. Vahdat, B. Y. Zhao, and H. Zheng. Mirror mirror on the ceiling: flexible wireless links for data centers. In *SIGCOMM 2012*.